

운영체제론 Project-2

FUSE-based File System Implementation

2019.05.23

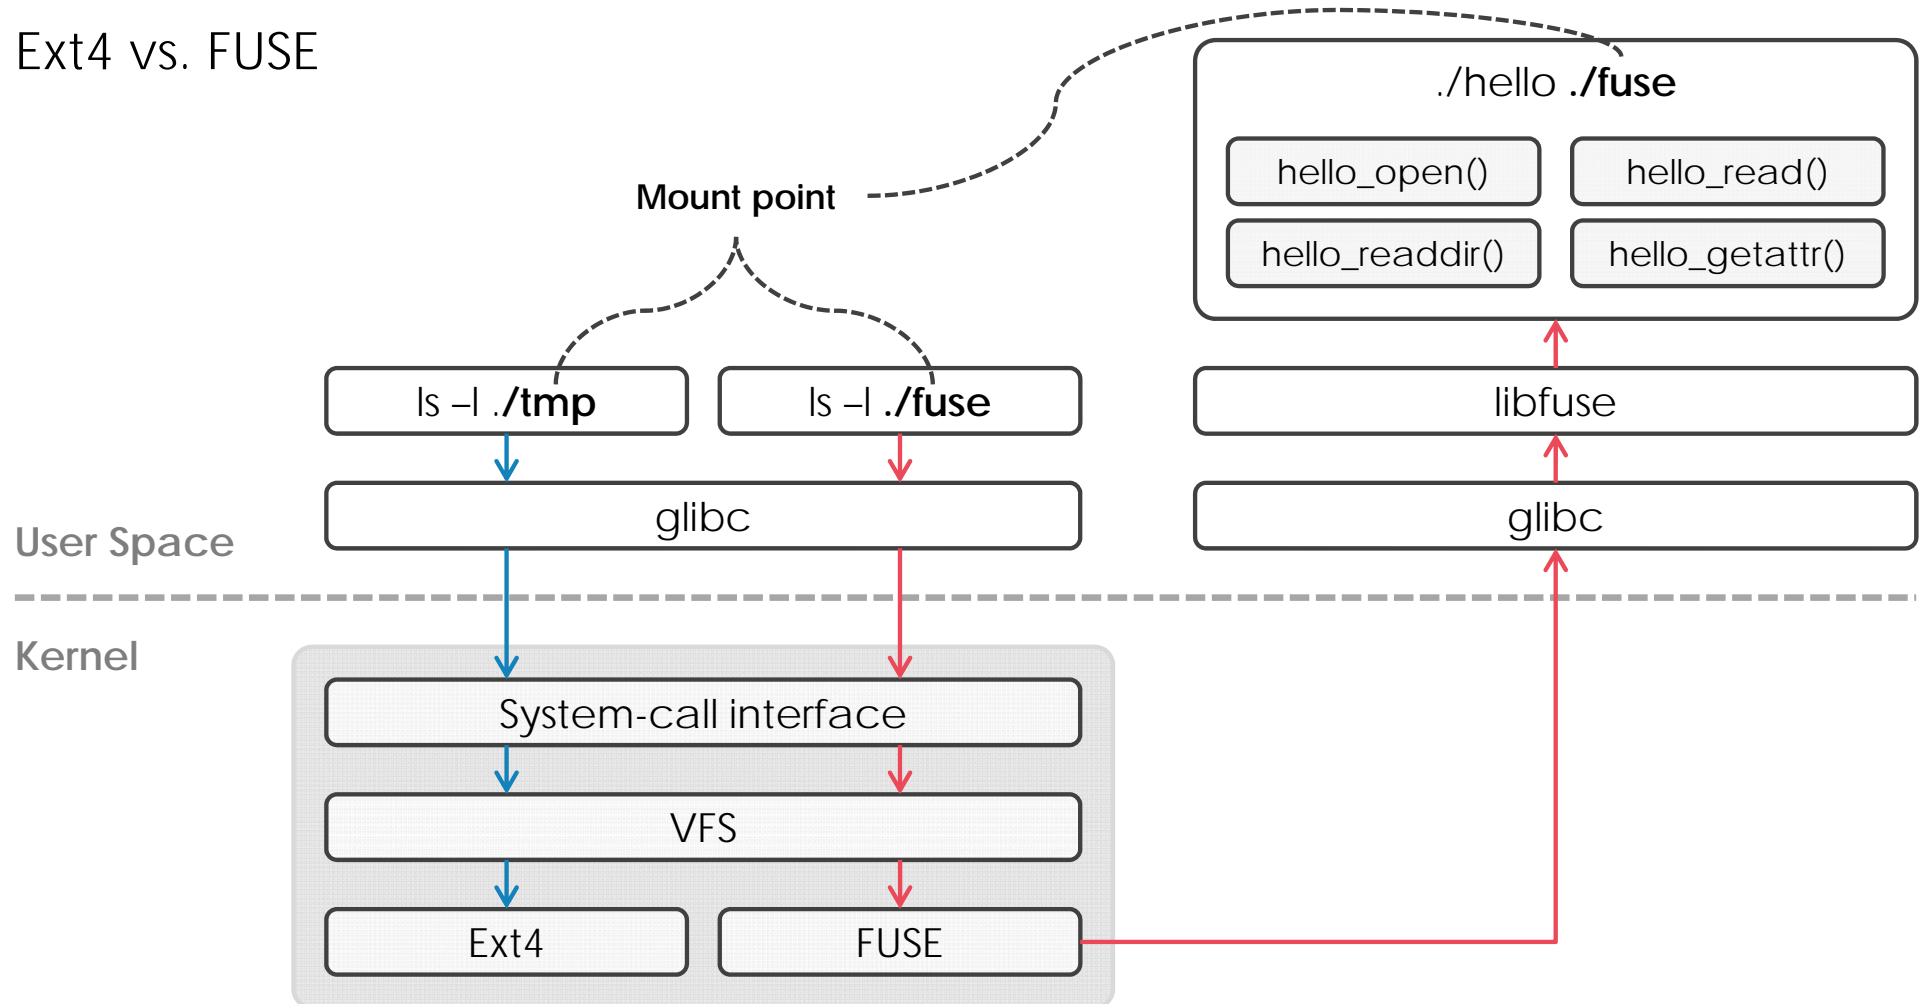
Contents

<u>02</u>	<u>03</u>	<u>04</u>	<u>05</u>	<u>06</u>
FUSE Introduction	FUSE Operations	Sample & Demo	Project	References

FUSE Introduction

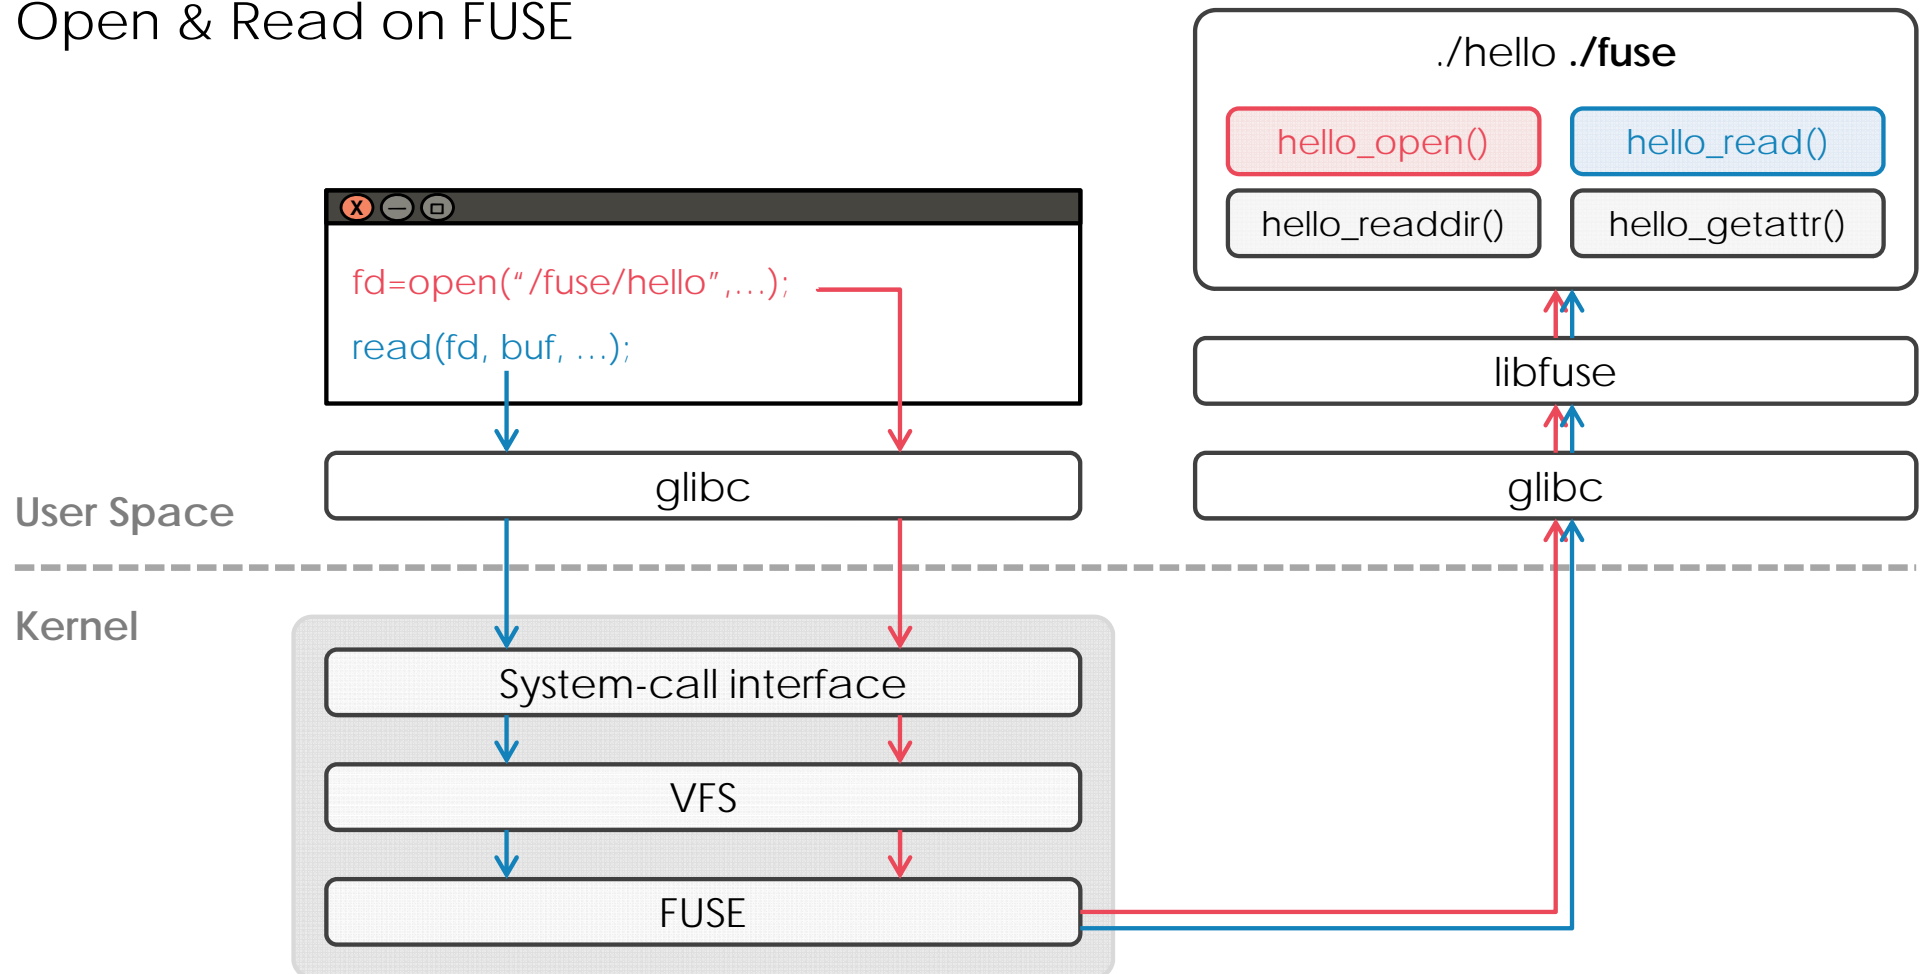
► FUSE (Filesystem in USErspace)

▪ Ext4 vs. FUSE



FUSE Introduction

- FUSE (Filesystem in USErspace)
 - Open & Read on FUSE



FUSE Operations

- ▶ File/Directory Management
 - File/Directory Descriptor (Tree)
- ▶ Fuse Operations
 - File Operations
 - Directory Operations
 - Metadata Operations
 - Other Operations
- ▶ Useful Errors in FUSE
- ▶ FUSE context

FUSE Operations

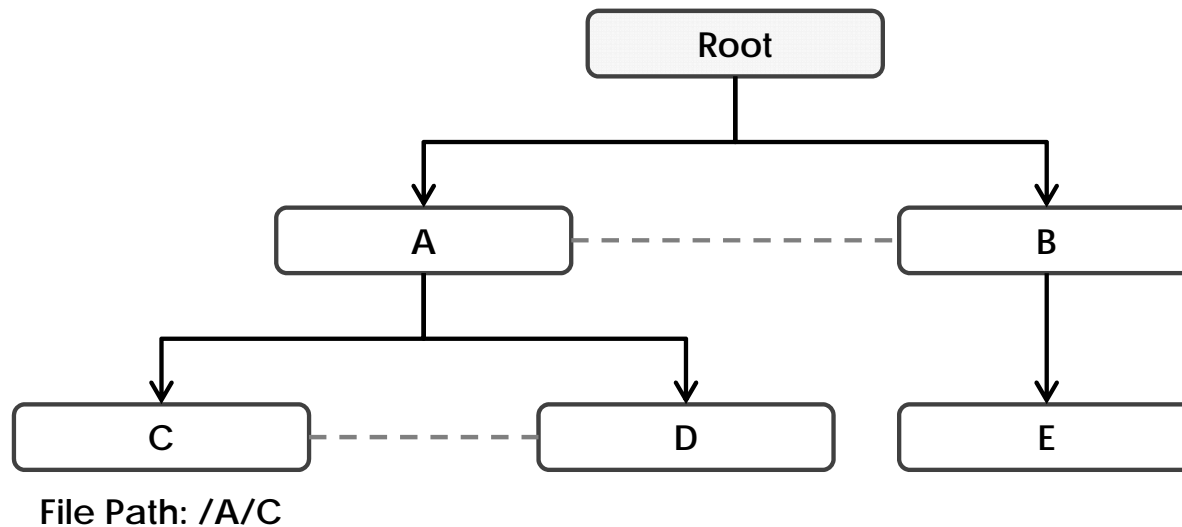
► File/Directory Descriptor

- size: size in bytes
- mode: type and permissions
- uid: owner id
- gid: group id
- atime: access time (often fudged)
- mtime: modification time
- ctime: metadata change time
- data: file data

FUSE Operations

► File/Directory Descriptor (Tree)

- next: next node pointer
- child: child node pointer
- parents: parents node pointer



FUSE Operations

► File Operations

- `mknod(path, mode, dev)`: create a file (or device)
- `unlink(path)`: remove (delete) the give file, link, etc.
- `rename(old, new)`: move and/or rename a file
- `open(path, flags)`: open a file
- `read(path, length, offset, fi)`
- `write(path, buf, offset, fi)`
- `truncate(path, len, fi)`: cut off at length
- `flush(path, fi)`: one handle is closed
- `release(path, fi)`: file handle is completely closed (no errors)

FUSE Operations

► Directory Operations

- `readdir(path)`: yield directory entries for each file in the directory
- `mkdir(path, mode)`: create a directory
- `rmdir(path)`: delete an empty directory

► Metadata Operations

- `getattr(path)`: read metadata
- `chmod(path, mode)`: alter permissions
- `chown(path, uid, gid)`: alter ownership

FUSE Operations

► Other Operations

- `statfs(path)`
- `fsdestroy()`
- `create(path, flags, mode)`
- `utimens(path, times)`
- `readlink(path)`
- `symlink(target, name)`
- `link(target, name)`
- `fsync(path, fdatsync, fi)`
- `etc.`

FUSE Operations

► fuse_operations [Link] https://libfuse.github.io/doxygen/structfuse__operations.html

int(* getattr)(const char *, struct stat *, struct fuse_file_info *fi)
int(* readlink)(const char *, char *, size_t)
int(* mknod)(const char *, mode_t, dev_t)
int(* mkdir)(const char *, mode_t)
int(* unlink)(const char *)
int(* rmdir)(const char *)
int(* symlink)(const char *, const char *)
int(* rename)(const char *, const char *, unsigned int flags)
int(* link)(const char *, const char *)
int(* chmod)(const char *, mode_t, struct fuse_file_info *fi)
int(* chown)(const char *, uid_t, gid_t, struct fuse_file_info *fi)
int(* truncate)(const char *, off_t, struct fuse_file_info *fi)
int(* open)(const char *, struct fuse_file_info *)
int(* read)(const char *, char *, size_t, off_t, struct fuse_file_info *)
int(* write)(const char *, const char *, size_t, off_t, struct fuse_file_info *)
int(* statfs)(const char *, struct statvfs *)
int(* flush)(const char *, struct fuse_file_info *)
int(* release)(const char *, struct fuse_file_info *)
int(* fsync)(const char *, int, struct fuse_file_info *)
int(* setxattr)(const char *, const char *, const char *, size_t, int)
int(* getxattr)(const char *, const char *, char *, size_t)
int(* listxattr)(const char *, char *, size_t)
int(* removexattr)(const char *, const char *)

FUSE Operations

► fuse_operations [Link] https://libfuse.github.io/doxygen/structfuse__operations.html

```
int(* opendir )(const char *, struct fuse_file_info *)
int(* readdir )(const char *, void *, fuse_fill_dir_t, off_t, struct fuse_file_info *, enum fuse_readdir_flags)
int(* releasedir )(const char *, struct fuse_file_info *)
int(* fsyncdir )(const char *, int, struct fuse_file_info *)
void (* init )(struct fuse_conn_info *conn, struct fuse_config *cfg)
void (* destroy )(void *private_data)
int(* access )(const char *, int)
int(* create )(const char *, mode_t, struct fuse_file_info *)
int(* lock )(const char *, struct fuse_file_info *, int cmd, struct flock *)
int(* utimens )(const char *, const struct timespec tv[2], struct fuse_file_info *fi)
int(* bmap )(const char *, size_t blocksize, uint64_t *idx)
int(* ioctl )(const char *, unsigned int cmd, void *arg, struct fuse_file_info *, unsigned int flags, void *data)
int(* poll )(const char *, struct fuse_file_info *, struct fuse_pollhandle *ph, unsigned *reventsp)
int(* write_buf )(const char *, struct fuse_bufvec *buf, off_t off, struct fuse_file_info *)
int(* read_buf )(const char *, struct fuse_bufvec **bufp, size_t size, off_t off, struct fuse_file_info *)
int(* flock )(const char *, struct fuse_file_info *, int op)
int(* fallocate )(const char *, int, off_t, off_t, struct fuse_file_info *)
ssize_t(* copy_file_range )(const char *path_in, struct fuse_file_info *fi_in, off_t offset_in, const char *path_out, struct fuse_file_info *fi_out, off_t offset_out, size_t size, int flags)
```

FUSE Operations

► Useful Errors in FUSE

■ Error Lists

- `errno.ENOSYS`: Function not implemented
- `errno.EROFS`: Read-only file system
- `errno.EPERM`: Operation not permitted
- `errno.EACCES`: Permission denied
- `errno.ENOENT`: No such file or directory
- `errno.EIO`: I/O error
- `errno.EEXIST`: File exists
- `errno.ENOTDIR`: Not a directory
- `errno.EISDIR`: Is a directory
- `errno.ENOTEMPTY`: Directory not empty

FUSE Operations

► Useful Errors in FUSE

- Ex) Error Check

→ `errno.ENOENT`: No such file or directory

```
static int hello_readdir(const char *path, void *buf, fuse_fill_dir_t filler,  
                        off_t offset, struct fuse_file_info *fi)  
{  
    (void) offset;  
    (void) fi;  
  
    if (strcmp(path, "/") != 0)  
        return -ENOENT;  
  
    filler(buf, ".", NULL, 0);  
    filler(buf, "..", NULL, 0);  
    filler(buf, hello_path + 1, NULL, 0);  
  
    return 0;  
}
```

FUSE Operations

► FUSE Context

- `fuse_get_context()` returns struct `fuse_context` with:
 - uid: User ID of the calling process
 - gid: Group ID of the calling process
 - pid: Thread ID of the calling process
 - `private_data`: Private filesystem data
 - umask: Umask of the calling process
- Useful for nonstandard permission models and other user-specific behavior

Sample

► Hello File System

- A single file (Hello) can be read (without write/create/delete)

[Link] <https://github.com/fuse4x/fuse/blob/master/example/hello.c>

- ex) hello.c - main function

```
#define FUSE_USE_VERSION 26

#include <fuse.h>
#include <stdio.h>
#include <string.h>
#include <errno.h>
#include <fcntl.h>

static const char *hello_str = "Hello World!\n";
static const char *hello_path = "/hello";

static struct fuse_operations hello_oper = {
    .getattr      = hello_getattr,
    .readdir      = hello_readdir,
    .open         = hello_open,
    .read         = hello_read,
};

int main(int argc, char *argv[])
{
    return fuse_main(argc, argv, &hello_oper, NULL);
}
```


Sample

► Hello File System

- ex) hello.c - Open function: Check file path and access rights to FILE

```
static int hello_open(const char *path, struct fuse_file_info *fi)
{
    if (strcmp(path, hello_path) != 0)
        return -ENOENT;

    if ((fi->flags & 3) != O_RDONLY)
        return -EACCES;

    return 0;
}
```

- ex) hello.c - Read function: Send corresponding data to file using file offset

```
static int hello_read(const char *path, char *buf, size_t size, off_t offset,
                     struct fuse_file_info *fi)
{
    size_t len;
    (void) fi;
    if (strcmp(path, hello_path) != 0)
        return -ENOENT;

    len = strlen(hello_str);
    if (offset < len) {
        if (offset + size > len)
            size = len - offset;
        memcpy(buf, hello_str + offset, size);
    } else
        size = 0;

    return size;
}
```

Sample

► Hello File System

- ex) hello.c - Readdir function: send the list of files in the directory to buf

```
static int hello_readdir(const char *path, void *buf, fuse_fill_dir_t filler,  
                        off_t offset, struct fuse_file_info *fi)  
{  
    (void) offset;  
    (void) fi;  
  
    if (strcmp(path, "/") != 0)  
        return -ENOENT;  
  
    filler(buf, ".", NULL, 0);  
    filler(buf, "..", NULL, 0);  
    filler(buf, hello_path + 1, NULL, 0);  
  
    return 0;  
}
```

Sample

► Hello File System

- ex) hello.c - Getattr function: send the attribute of path file

```
static int hello_getattr(const char *path, struct stat *stbuf)
{
    int res = 0;

    memset(stbuf, 0, sizeof(struct stat));
    if (strcmp(path, "/") == 0) {
        stbuf->st_mode = S_IFDIR | 0755;
        stbuf->st_nlink = 2;
    } else if (strcmp(path, hello_path) == 0) {
        stbuf->st_mode = S_IFREG | 0444;
        stbuf->st_nlink = 1;
        stbuf->st_size = strlen(hello_str);
    } else
        res = -ENOENT;

    return res;
}
```

Demo

► Install the FUSE library on Ubuntu

```
sudo apt-get install libfuse-dev
```

- [Link] <https://github.com/libfuse/libfuse/>

► Source code download (Optional)

- [Link] <https://github.com/fuse4x/fuse>
- [Link] <https://github.com/fuse4x/fuse/blob/master/example/hello.c>

```
minhozx@dclab-PowerEdge-R940:~/Workspaces$ git clone https://github.com/fuse4x/fuse
Cloning into 'fuse'...
remote: Enumerating objects: 8129, done.
remote: Total 8129 (delta 0), reused 0 (delta 0), pack-reused 8129
Receiving objects: 100% (8129/8129), 2.11 MiB | 553.00 KiB/s, done.
Resolving deltas: 100% (6388/6388), done.
Checking connectivity... done.
minhozx@dclab-PowerEdge-R940:~/Workspaces$ ls
fuse
minhozx@dclab-PowerEdge-R940:~/Workspaces$ cp fuse/example/hello.c ./
minhozx@dclab-PowerEdge-R940:~/Workspaces$ ls
fuse hello.c
minhozx@dclab-PowerEdge-R940:~/Workspaces$
```

Demo

► Compile

```
gcc -D_FILE_OFFSET_BITS=64 -o hello hello.c -lfuse
```

```
minhozx@dclab-PowerEdge-R940:~/Workspace$ gcc -D_FILE_OFFSET_BITS=64 -o hello hello.c -lfuse
minhozx@dclab-PowerEdge-R940:~/Workspace$ ls
fuse hello hello.c
minhozx@dclab-PowerEdge-R940:~/Workspace$
```

► Mount

- ./hello [mount point]

```
./hello ./mnt
```

```
minhozx@dclab-PowerEdge-R940:~/Workspace$ gcc -D_FILE_OFFSET_BITS=64 -o hello hello.c -lfuse
minhozx@dclab-PowerEdge-R940:~/Workspace$ ls
fuse hello hello.c
minhozx@dclab-PowerEdge-R940:~/Workspace$ mkdir mnt
minhozx@dclab-PowerEdge-R940:~/Workspace$ ./hello ./mnt/
minhozx@dclab-PowerEdge-R940:~/Workspace$ mount | grep 'mnt'
/home/minhozx/Workspace/hello on /home/minhozx/Workspace/mnt type fuse.hello (rw,nosuid,nodev,relatime,user_id=1003,group_id=1003)
minhozx@dclab-PowerEdge-R940:~/Workspace$
```

Demo

► Reading a file

- `cat [file]`

```
cat hello
```

```
minhozx@dcclab-PowerEdge-R940:~/Workspace$ mount | grep 'mnt'
/home/minhozx/Workspace/hello on /home/minhozx/Workspace/mnt type fuse.hello (rw,nosuid,nodev,relatime,user_id=1003,group_id=1003)
minhozx@dcclab-PowerEdge-R940:~/Workspace$ cd mnt/
minhozx@dcclab-PowerEdge-R940:~/Workspace/mnt$ ls -al
total 4
drwxr-xr-x 2 root    root      0 1월  1 1970 .
drwxrwxr-x 4 minhozx minhozx 4096 5월 22 23:14 ..
-r--r--r-- 1 root    root      13 1월  1 1970 hello
minhozx@dcclab-PowerEdge-R940:~/Workspace/mnt$ cat hello
Hello World!
minhozx@dcclab-PowerEdge-R940:~/Workspace/mnt$
```

Demo

► Unmount

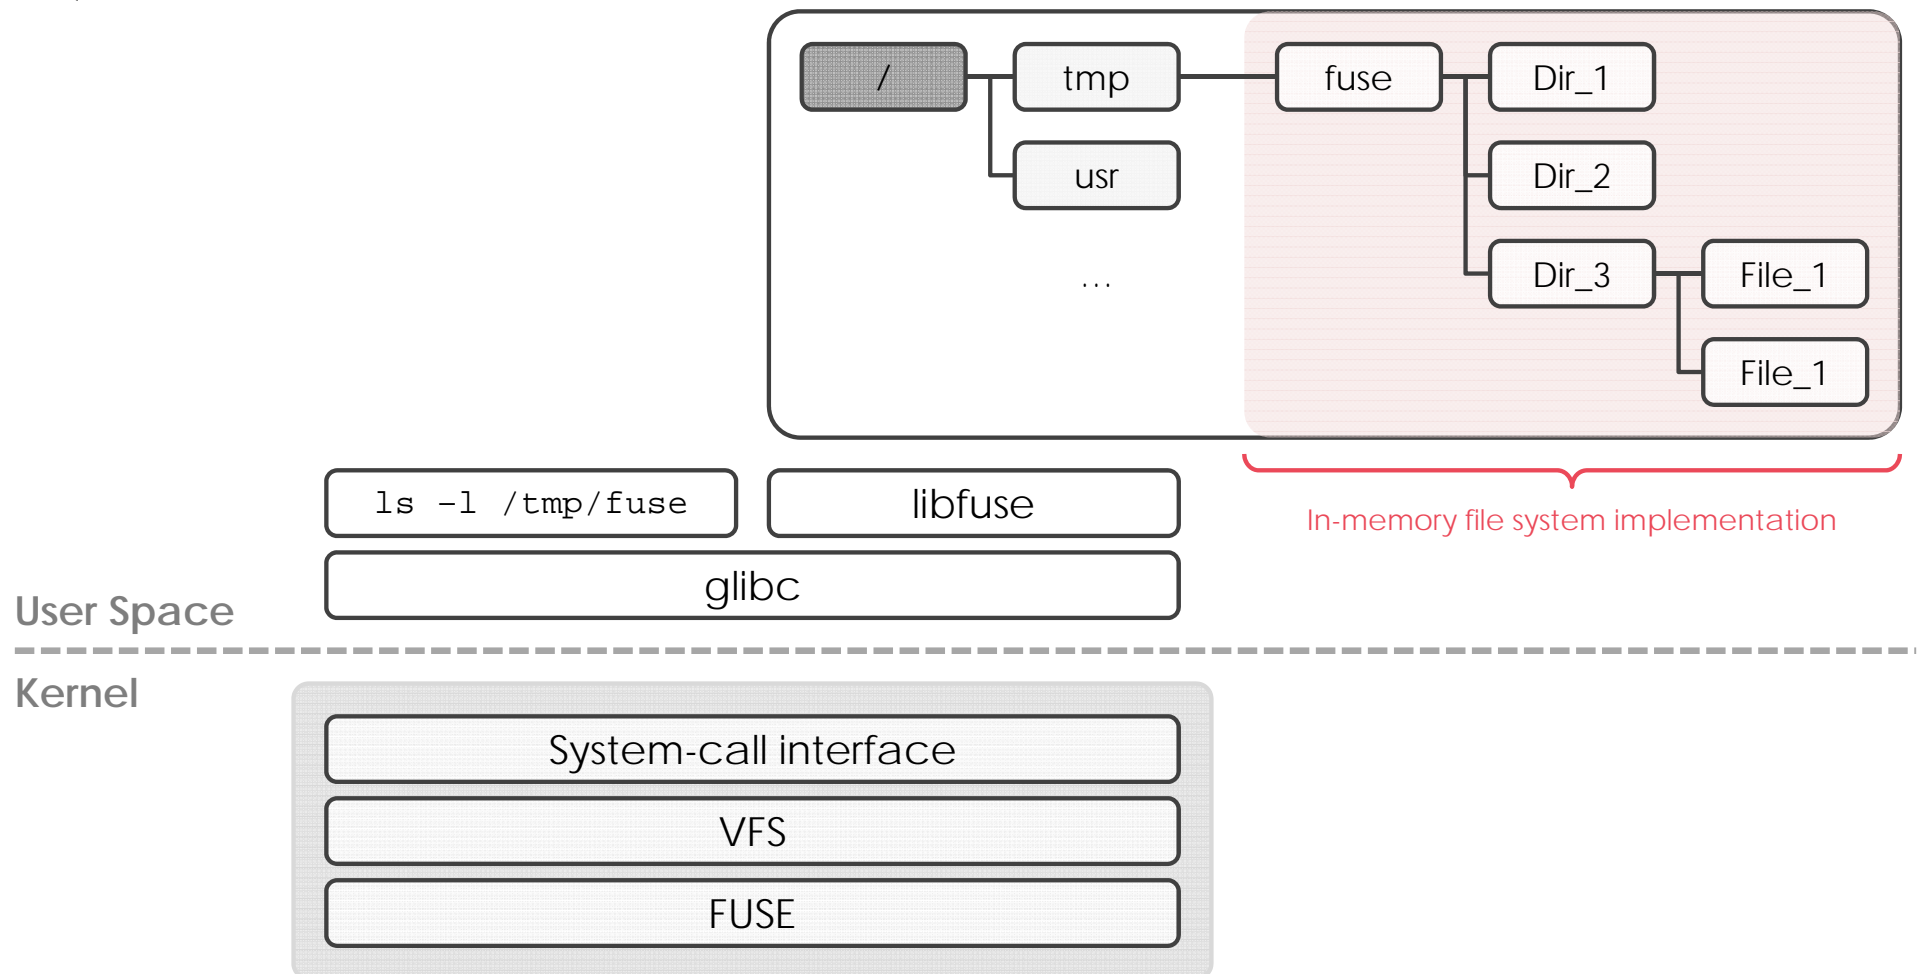
- `fusermount -u [mount point]`

```
fusermount -u ./mnt
```

```
minhozx@dcclab-PowerEdge-R940:~/Workspace$ mount | grep 'mnt'
/home/minhozx/Workspace/hello on /home/minhozx/Workspace/mnt type fuse.hello (rw,nosuid,nodev,relatime,user_id=1003,group_id=1003)
minhozx@dcclab-PowerEdge-R940:~/Workspace$ cd mnt/
minhozx@dcclab-PowerEdge-R940:~/Workspace/mnt$ ls -al
total 4
drwxr-xr-x 2 root    root      0 1월  1 1970 .
drwxrwxr-x 4 minhozx minhozx 4096 5월 22 23:14 ..
-r--r--r-- 1 root    root      13 1월  1 1970 hello
minhozx@dcclab-PowerEdge-R940:~/Workspace/mnt$ cat hello
Hello World!
minhozx@dcclab-PowerEdge-R940:~/Workspace/mnt$
```

Project

- ▶ FUSE-based in-memory file system implementation
 - ex) tree structure



Project

▶ FUSE-based in-memory file system implementation

■ 필수 기능 구현 [기본]

- 파일 열기/닫기 (open/close)
- 파일 읽기/쓰기 (read/write)
- 파일 생성/삭제 (create/remove)
- File descriptor를 이용한 기능 구현
- 다단계 디렉터리 생성 가능 ex) A/B/C/D/file
- 상황에 맞는 에러코드(errno) 3개 이상 사용

■ 추가 기능 구현 [가산점]

- 파일의 r/w/x 권한 변경
- 파일 이동/복사 (move/copy)
- 파일 링크 (link) [softlink/hardlink 구분]
- 기타 각자 원하는 추가 기능을 구현 가능

■ 주의사항

- 유저 터미널로의 직접 접근 금지 ex) printf("error = ... ")
- Error를 나타낼 때 'return -error code' 사용
- 필수 및 추가 기능에 대한 설명은 보고서에 '반드시' 명시할 것

Project

▶ 수행방법

- 2인 1팀으로 팀 편성
- 팀 단위 과제 수행 및 결과물 제출

▶ 제출물

- 1차 - 제안서
 - FUSE 조사/분석서
 - In-memory File System 설계서
(자료구조 및 파일 데이터 관리 방식, 필요한 함수에 대한 설명 등)
- 2차 - 결과보고서
 - 구현 기능 설명서
 - 실행 결과 화면 설명서 (스크린샷 첨부)
 - 소스 코드

Project

▶ 제출기한

■ 제안서 제출

→ 2019.06.04(화) 오전 11:00 까지

→ Hardcopy로 출력하여 27315호(분산컴퓨팅연구실)에 제출

■ 결과보고서 제출

→ 2017.06.13(목) 오전 11:00 까지

→ 결과보고서와 소스 코드를 압축하여 iCampus 과제란에 제출

References

- ▶ FUSE Official Webpage, <https://github.com/libfuse/libfuse>
- ▶ FUSE Document – fuse_operations list,
http://fuse.sourceforge.net/doxygen/structfuse__operations.html
- ▶ https://www.cs.hmc.edu/~geoff/classes/hmc.cs135.201001/homework/fuse/fuse_doc.html
- ▶ M. Abd-El-Malek, M. Wachs, J. Cipar Gregory, R. Ganger, G. A. Gibson¹, M. K. Reiter, File system virtual appliances: Third-party file system implementations without the pain, Carnegie Mellon University Panasas, Inc., 2008